

Divinci AI — Investor Executive Summary: companion, v0.6

Prepared by Michael Mooring, Samuel Tobia and Anindya Sundar Nandi. 25 September 2026. The detail behind Divinci AI's one-page Investor Executive Summary v0.6. Changes from v0.5: our own September re-checks weakened two results we had reported, and the person-level test has now run; both are reported below as they stand. The money section now reconciles the burn, break-even and programme figures, and the gate adds the evidence the SBDC review asked for.

1. Company overview

Name and founding date.

Divinci AI Inc., a Delaware C-corporation (file no. 2970809), incorporated 22 January 2024. We started building the platform in mid-2023. Based in Irvine, California. Website: divinci.ai.

Mission.

Providing AI safety, security, and governance in an AI-native form factor.

What we offer, in plain terms.

The Divinci platform builds and runs custom AI assistants for businesses — assistants that answer from a company's own documents, work on its website, and plug into its tools. The product this summary is about is our compliance product, Model Erasure with a receipt. It finds what an AI model has learned about a person or a work, removes it from the model's own memory, checks that the model still works, and hands the auditor a signed receipt they can verify on their own, without having to take our word for anything.

The platform is the foundation the compliance product runs on.

2. Problem

What pain are we solving?

A company can delete a person's data from its database. But an AI model that learned from that data can still remember the person. The database gets cleaned; the memory doesn't.

In Europe that is a legal problem twice over.

Privacy. GDPR, in force since 2018, gives people the right to have their data erased. France's privacy regulator, CNIL, said in February 2025 that by default, honoring that right for an AI model means retraining the model, and that retraining can take up to three months. Retraining is slow and expensive, and is rarely done. We have not found a large European AI operator with a tested answer ready for a

demand that data be removed from a model.

Copyright. Munich Regional Court I ruled on 11 November 2025 (GEMA v. OpenAI, 42 O 14139/24) and again on 31 July 2026 (GEMA v. Suno) that a model which has memorised a work reproduces it, and that offering such a model to the public is itself an act of communication to the public — separate from the training, and separate from anything a user does with it afterwards. If that reading holds, releasing model weights that still carry a work is the exposure, and the only place to fix it is the weights before they go out. GEMA is Germany's music collecting society, and its counterparts across Europe have the same interest in those judgments. This is a second buyer with the same technical need and a much sharper clock: a release date.

In the United States the same question has just been written into law, in California. Since 1 January 2025 the California Consumer Privacy Act's definition of personal information expressly includes information held in "artificial intelligence systems that are capable of outputting personal information" (AB 1008, Civil Code § 1798.140(v)). The state's privacy regulator, the California Privacy Protection Agency, wrote in its own July 2024 staff memo that this makes such information subject to deletion requests and to data minimisation. The Attorney General's January 2025 legal advisory says the same. Nobody has been fined on it yet, and we have found no regulator or vendor guidance on how a deletion request reaches a model. That gap is our product.

One thing we say plainly everywhere: no regulator, in Europe or in California, has yet ruled that any removal method fully satisfies the erasure right. That question is open for the whole market. So we sell a documented, measured, checkable way to respond — not a legal guarantee. CNIL itself has written that requests a company can refuse today may have to be honored tomorrow as removal techniques improve. We're building for that tomorrow.

How big and how urgent?

Fines under GDPR run up to €20 million or 4% of worldwide revenue, whichever is higher. California fines run up to \$7,500 per intentional violation, per consumer, and the state has collected settlements from \$1.2 million (Sephora, 2022) to \$12.75 million (General Motors, 2026) — so far for opt-out, children's-data and data-broker failures; deletion appears only as a remedy in those orders, never yet as the violation. The alternative we replace, full retraining, can cost anywhere from tens of thousands to several hundred thousand dollars per model once people, data preparation, engineering time and downtime are counted — and it takes months.

One timing note: Europe's AI Act did not switch on its strict rules this August as originally scheduled. A July 2026 law, the Digital Omnibus, pushed those rules out to December 2027 and August 2028. So we no longer lean on the AI Act for urgency. GDPR has been live for eight years, the Munich copyright rulings are live now, and California's statute has been live since January 2025.

3. Solution and benefits

How the product works.

Three steps: find, remove, prove. We find what the model learned about a person or a work. We remove it. For data a customer fine-tuned in, that means today rebuilding the fine-tune without the person from a verified public base, never the base model itself; cheaper routes are in testing and not yet claimed. Then we prove two things — that it's gone, and that the model still works — by measuring the model before and after against quality bars set in advance. The result is a receipt carrying a standard digital signature, published at a public address, built so a regulator, auditor or customer can check it without needing anything from us. The checking tool is written and open-licensed; publishing it is part of the 60-day plan in section 5.

This works on models a company downloads and runs itself: Llama, Mistral, Qwen, Gemma and their relatives. It doesn't apply to closed services like ChatGPT, which only their makers can change.

What makes us different.

The proof step is the product. The editing science is published; anyone can use it, and eventually competitors will. What nobody sells today — and we checked this September, properly, in Europe and again in California — is the whole bundle: measure before, remove, measure after, and hand over signed evidence a stranger can verify. The leading removal company reports its own before-and-after test scores and asks customers to trust them. The big governance platforms produce paperwork about models; they don't change models. The privacy-software companies delete from databases and, at best, from training datasets — Transcend's "Deep Deletion" and BigID both say in their own material that they stop before the model. Provable removal exists in academic papers, not on the market.

Two claims we deliberately never make. We don't say we invented unlearning — an Israeli company called Hirundo has claimed "first" since 2023, and they're real. And we don't pretend the receipt is hard technology; a digital signature is ordinary. Our claim is narrower: as of September 2026 we have found no one selling the whole bundle — measure, remove, measure again, and hand over signed evidence a stranger can verify.

What we have actually shown, told honestly.

Fact-level editing of a model's own weights: published, and since qualified by our own re-checks.

On 25 April 2026 we published an edit that suppressed one fact — "Paris is the capital of France" — in Google's Gemma 4 E2B model by editing one feature in one layer (feature 11179, layer 27). The internal signal carrying the fact fell from 18.10 to below the model's 25 strongest signals, and overall language quality moved by +0.02% on WikiText-103. The change is a checksummed patch of about 9 KB, and removing it restores the original model byte for byte. The April post checked four other prompts; two later findings of ours say that is not enough, and we report them here. First, on 17 September a stronger deletion of the same fact (512 neurons chosen by attribution) came back after two small retraining steps on one public sentence: for facts learned in pretraining, what we can do today is suppression, not removal. Second, on 24 September a wider check of an edit made with our shipping toolchain (installing a made-up fact into the same model on 20 September) found that, although the file changed in only one

row and restores byte for byte, the edit damaged 8 of 100 other facts, about one neighbouring capital in five, and leaked its inserted answer into other questions. So we no longer describe edits of this kind as surgical until they pass the wider check. We make no erasure claim for pretraining data. Method: divinci.ai/blog/deleting-paris-from-a-language-model.

The August test shows the measurement discipline. That test removed documents from a model's reference library, not from its weights, and involved no personal data. Our measurements caught the untouched control topics drifting more than the erased one (-0.058 against -0.026) — the kind of thing a vendor who doesn't measure would ship without ever knowing. It is documented in full in our engineering record, negative result included.

Person-level deletion: the test has run and has not yet passed. A fact is one feature; a person is many. In September we fine-tuned records about five made-up people into Gemma 4 E2B, removed them by rebuilding the fine-tune without them from the verified public base, and compared the result with six control models that never saw those people. Five pass/fail checks were written down before the run: the base is the publisher's genuine file (pass); the base is untouched afterwards (pass); the result can be re-derived by an independent rebuild (pass); the people cannot be recovered by any informative attack (9 of 9, pass); and the model is still usable (undecided). Instruction-following scored $0.931\times$ the base, with an uncertainty interval of $0.862\text{--}0.983$ that straddles the 0.90 bar. The rule set before the run says the recovery result counts as erasure evidence only once usability passes, so the test has not passed. The model that learned the people sits on the same margin, which points at the fine-tuning recipe rather than the removal; the bar does not move. Cheaper methods than a rebuild did not meet their bars: gradient unlearning either missed the erasure bar or cost retained knowledge, and weight edits do not remove a fine-tuned-in person. It has run on one small model family only. Passing it, and repeating it on a larger model from a second family, is the first item of the 30-day gate.

What the customer gets.

An erasure demand answered with evidence: the rebuild, or a cheaper method once one passes, measured before and after. Evidence an auditor can check alone. Damage to the model measured, not assumed. One artifact that a regulator, an auditor, a customer — and increasingly an insurer — can verify independently.

4. Target market

Who buys.

The people on the hook when an erasure demand arrives: privacy officers, compliance lawyers, and the risk teams at organisations running downloadable models with European or Californian users. Four groups: compliance lawyers, data-protection officers, compliance insurers, and the software firms that implement compliance. We reach them directly, through our seven consultancy partners, and — early conversations only — through insurers who cover AI risk. We opened that insurer outreach this month, to a seven-company list that includes Armilla, Munich Re and Vouch; the first call, with Vouch, was scheduled for 14 September.

Two further segments. First, *open-weight model publishers* — the companies that release model weights publicly (in Europe: Mistral, Aleph Alpha, Black Forest Labs, Kyutai, H Company, Stability AI). After the Munich rulings, the release itself is their exposure, and the buyer is the general counsel; we have verified named legal contacts at five of them this week, and have contacted none of them yet. Second, *California-exposed businesses* under the amended CCPA: registered data brokers (now under a state deletion platform charging \$200 per unanswered request per day), ad-tech and identity-graph companies, health and wellness apps, and consumer AI companies. Here the buyer is the chief privacy officer, and the channel is the auditor: the state's new cybersecurity-audit rule (§ 7122(d)) requires audit findings to rest on evidence the auditor can check rather than on management's assertions. Whether a signed erasure receipt satisfies that requirement is untested; it is built to.

The insurance angle matters more than it might look. Armilla already sells AI liability insurance backed by Lloyd's of London, and its policy includes an independent certification of the AI system being insured, and coverage limits passed \$25 million per company in January. Munich Re and HSB have signalled that they favour documented AI governance. Nobody yet requires a signed erasure receipt to underwrite a policy. Our receipt is a candidate to become exactly that, and that's the conversation we're opening.

How big is the market?

Honestly: small and unmeasured today — likely tens to low hundreds of qualified organisations worldwide, larger now that California and copyright are in scope, but still unmeasured. Measuring it properly, through 15 to 20 interviews with privacy lawyers and data-protection officers, is one of the 30-day tasks, not a number we assume. The gate also produces a bottom-up count: organisations that publish or fine-tune open-weight models and have EU or California users, built from public model registries, EU company filings and the partners' own client lists, each with a named buyer role. Until that count exists we quote no market size of our own. For context only, an unattributed industry projection puts the machine-unlearning niche growing from about \$0.15 billion in 2025 to about \$2.7 billion by 2032, fast growth off a very small base; we treat that as background, not as our plan. We could not verify the "Fortune 500 compliance budget" figure in circulation, so it does not appear here.

5. Traction and milestones

Built and working today.

Our receipt system is live in production: the first real certificate has been issued, its signature is valid, and anyone can verify it with zero credentials. The public key registry is live. Our quality-testing system is live, with automatic pass/fail bars, scheduled runs and alerts. The erasure tool works on downloadable models, with two API surfaces, two developer kits, a command-line tool and a web interface; the fact-level deletion in section 3 was produced with it. We've published nine model indexes and catalogued sixteen. Seven software consultancies have signed partnerships and are ready to run our demo pipeline. We have a product demo video, membership in the Google for Startups Cloud Program with a Google Cloud Marketplace listing in progress, and a registered trademark. We are also a partner under Cloudflare's Workers Launchpad accelerator program, and a member of JPMorgan's Innovation Economy programs.

Since the last version. The European lawyer is scoped: the Hamburg office of an international firm with data-protection and copyright practices. On 14 September we agreed in principle a customer-facing memo plus an internal note on the gray zones, with fees capped at €20k. The opinion rests on an adversarial test designed and run by an external German AI engineering firm, for about €10k more. Nothing is instructed yet; a follow-up call with our investor and advisor is being arranged to settle deliverables and the test partner. An existing investor plans to cover the ~€30k. We executed our first vendor HIPAA business-associate agreement (with OpenAI, 11 September); our HIPAA-scoped workspaces refuse to send data to a model provider not under such an agreement, the same fail-closed discipline the receipt is built on. We met the Mucker Accelerator on 14 September following our application.

What has not started. Buyer interviews, design-partner agreements and partner-sourced prospects are all at zero. The seven consultancy partnerships are signed, but none has yet produced an introduction, meeting or pilot. We state that here because the gate measures it.

Still to prove: the 30-day plan.

The gate runs for 30 days. Each line has its pass/fail bar written down before it runs, so the result means something either way.

1. **Person-level erasure.** All five checks pass on Gemma 4 E2B, then on a mid-sized model from a second family (target: DeepSeek-V4.1-Flash), against control models that never saw the person. Any failed or undecided check is a fail, and we publish it.
2. **Independent validation.** The external German AI engineering firm designs and runs an adversarial test, trying to recover the person from the model by several routes; the lawyer's opinion on GDPR Article 17 and AI Act Article 53(1)(c) rests on it.
3. **Willingness to pay.** 15 to 20 interviews with privacy officers, DPOs and heads of AI. Each answers the same questions: who holds the budget, what event triggers a purchase, what they do today when an erasure request touches a model, and whether \$60–150k a year per model is within reach. We report the distribution, not the best quote.
4. **Design partners.** Three signed, each with a written scope (one model, one real or realistic erasure request), success criteria agreed in advance (the five checks, plus the partner's own quality bar), and a paid next step: if the criteria are met, a first-year readiness contract at list price less the design-partner discount.
5. **Partner activity, measured.** For each of the seven consultancies: introductions made, meetings held, qualified prospects (a named buyer, a model in scope, a stated erasure exposure). A partner that produces none in 30 days is reported as such.
6. **Cost per erasure.** Today we know only compute time: 3.4–6.7 GPU-hours on a single cloud GPU per test run, on the small model. The gate measures the full cost of one erasure on the second model: compute, engineering time and the re-measurement.
7. **Publication.** The evaluation method published within 60 days, failures included; the receipt-checking tool (already written, MIT-licensed, no dependency on Divinci) ships publicly with it.

If the plan passes its bar, one contingent step follows: an application to In-Q-Tel's technology development program (typically \$500k–3M, non-dilutive), led by the receipt. Being straight about it: we

have no relationship with In-Q-Tel today and no clearances. It's an option the gate result would unlock, nothing more.

6. Business model

How we charge.

We price below typical retraining costs, which run from tens of thousands to several hundred thousand dollars per model:

What	Price
One-time model mapping, per model version	\$30k–75k
Erasure readiness, per model per year (a set number of removals, with receipts)	\$60k–150k
Extra removals beyond the included number	\$2k–5k each
Partner revenue share	20–25% of list
Design partners	Free, in exchange for a dated, publishable case study

Can it scale and make money?

Revenue grows with the customer's list of models, not with our headcount, because the measuring and receipt steps are automated. Revenue today is zero. The first dollars should come from the consultancy partnerships and from design partners converting into paid contracts. We don't claim a profit margin yet, because we haven't measured what one job costs us to run; all we have is compute time for test runs on a small model (3.4–6.7 GPU-hours each). Measuring that is one of our first tasks, and we would rather say "unmeasured" than invent a percentage.

7. Team

Who we are.

Michael Duane Mooring, co-founder and CEO: built the Divinci platform from its first line of code in 2023 through everything now in production; leads product, engineering and this compliance push. Samuel Tobia, CTO: technical co-founder; infrastructure, deployments, and our Cloudflare and Google Cloud estate. Anindya Sundar Nandi, product marketing: demo pipeline, partner enablement and customer-facing evaluation. Two technical co-founders and five shipping product surfaces (web, embed, desktop, developer kit, MCP).

Advisors and partnerships.

The Brink SBDC @ USD (Bobby Edelman). The Google for Startups Cloud Program. Cloudflare's Workers Launchpad accelerator. JPMorgan's Innovation Economy programs. Seven consultancy partnerships. A named academic willing to attack and co-sign our erasure results is a recruitment target of the 30-day plan; we don't have one yet, and we'd rather say so than imply otherwise.

8. Financial highlights

Where the money stands.

Revenue is \$0. We're bootstrapped; no seed round has been raised. The numbers below replace v0.5's, which predated our Cloudflare Enterprise contract and understated the burn.

Burn today, founders unpaid	\$11,500–13,000 a month: Cloudflare Enterprise \$5,278.90, Google Cloud \$1,500–3,000 (startup credit now spent), other operating costs about \$4,700.
Break-even	\$17,500–19,000 a month: the burn plus \$2,000 a month for each of three founders.
Compliance programme, year 1	About \$224,000 (\$179,000 plus a 25% contingency), or about \$18,700 a month. This sits on top of the burn; it is not part of break-even.
Total year-1 cash need at \$0 revenue	Roughly \$362,000–380,000: twelve months of burn (\$138–156k) plus the programme (\$224k). Year one is funded initially by a seed/angel round, with a Series A planned for 2027.

Looking forward.

v0.5 gave break-even as about \$11,000 a month; with the Enterprise contract it is \$17,500–19,000, and the programme budget sits on top of it. We plan for zero programme revenue in year one. Our workbook holds three scenarios: Lean at about \$1.53M (five people, prove the product sells), Standard at about \$3.47M (eight or nine people, \$150–400k revenue from two to four direct deals), and Aggressive at about \$6.06M. Our plan is the Lean scenario, which is the seed/angel round, raised on the Divinci platform itself, whether or not the gate passes.

9. Funding ask

How much, and when.

Year one is funded initially by a seed/angel round, with a Series A planned for 2027. Now: a friends-and-family round. An existing investor is participating and plans to cover the gate's one outside cost, the legal opinion and the independent technical test (about €30k). The seed/angel round is about \$1.53M against the Lean plan, raised on the Divinci platform itself, whether or not the gate passes.

Two honest additions. The friends-and-family round is small and sized to the gate; the ~\$1.53M

seed/angel round does not wait on it. And what we're asking established investors for today isn't money. It's three things: design partners (three companies running downloadable models with European or Californian exposure), an introduction to California privacy counsel to match the European one now scoped, and that academic co-signer.

What the \$1.53M would buy.

Mostly people and proof. Engineering \$810k (two machine-learning engineers and one platform engineer). Independent validation plus the SOC 2 and ISO 42001 certifications big companies require, \$180k. Legal and regulatory \$90k. Computing and cloud \$55k. Design-partner delivery, events and travel \$65k. Insurance and overhead \$25k. A 25% contingency, \$306k. Total: \$1,225k plus 25%, or \$1,531k.

Return and valuation.

No valuation target is set. The honest return frame today is the break-even line in section 8.

10. Competitive advantage

Who we're really up against.

One company does what we do at the core: Hirundo, from Tel Aviv, founded in 2023 by a former Technion computer-science dean and two co-founders, with an \$8M seed round raised in June 2025. They remove learned data from the same model families we target. They're SOC 2 certified, they deploy anywhere including air-gapped environments, they joined HPE's AI partner program in March, Google DeepMind featured their work in May, and they run pilots in finance, healthcare and defense. On money, engineering maturity, certifications and distribution, they are ahead of us, and we say so.

Here's what they don't offer: evidence someone else can check. Their verification is their own before-and-after test scores, in their own reports — trust-us evidence. An independent analysis this May said exactly that: their results come from their own testing and haven't been outside-verified. We have not seen them publish a failure, and their marketing numbers are theirs, from their own tests, so we don't repeat them as fact. Selling that missing proof is our whole business.

Why an independent receipt matters to the buyer, not just to us. The person who pays for erasure is rarely the person who has to defend it. A DPO answering a regulator, an auditor signing off under California's audit rule (§ 7122(d), which requires findings to rest on evidence the auditor can check rather than management's assertions), or an insurer pricing AI liability all need evidence that does not depend on the vendor's own word. A vendor's self-reported test score gives them nothing they can check. A signed receipt with the before-and-after measurements, checkable with an open-source tool, does. If Hirundo adds receipts, the question becomes whose measurements are independently tested and whose failures are published; that is why the gate puts an outside adversarial test ahead of any sales push.

Everyone else is in a different lane. The governance platforms — Gartner published its first ranking of that category in June, with IBM, ServiceNow and others at the top — make policy documents, checklists

and sign-off records about models; none of them change a model. The privacy-software companies (BigID, Transcend, DataGrail and the like) delete from databases, apps and, in Transcend's case, training datasets — and say themselves that they stop before the model; for us they're future partners, not rivals, and Transcend is the first partner conversation we would open in California. And provable removal remains an academic subject. No one ships it.

What's actually hard to copy.

Not the signature and not the editing; both are ordinary. Three things compound over time. First, a growing pile of receipts that outsiders have actually checked, plus our record of publishing failures — a reputation no one can buy quickly, and no competitor has attempted. Second, our measurement method: what to test, how to set the bar beforehand, how to show nothing else broke. Third, the prize if we win it: regulators, auditors or insurers accepting our receipt as the format. California's audit rule, which requires findings to rest on checkable evidence rather than management's assertions, is the first regulatory text we know of that points at this kind of evidence. That last one is what would turn a product into a standard.

The threat we watch.

Hirundo adding a signed receipt. They already run the tests; if they chose to add receipts, we estimate the signing work in weeks. That's the clock we're racing, and it's why the 30-day plan front-loads the proof: a person-level result, a lawyer's opinion, a real case handled, and everything published. Behind that: a big governance platform buying its way into removal, our own erasure engine sitting in testing while rivals ship, a regulator ruling that removal can never satisfy the law (which would hurt the compliance pitch but make expensive retraining mandatory, raising the value of any cheaper path), and free academic tooling maturing.

Patents and property.

The editing methods are published science, so we own no secret there. Our tooling builds on an Apache-licensed open-source project, and our additions to it have been public since late August. We hold a registered trademark, have filed no patents, and have budgeted a freedom-to-operate review; the European lawyer's scope deliberately excludes patents, which need US patent counsel. The barrier is the running evidence infrastructure, the method, and the track record — not filings.

The six questions investors will press on

Will regulators accept verified erasure instead of retraining?

Nobody knows yet. Retraining is the stated default in Europe and no ruling exists either way; California has written the obligation into law without saying how it is met. We sell risk reduction and documentation today, and the lawyer's opinion inside the 30-day plan is our first formal read. CNIL's own guidance points toward removal techniques becoming acceptable as they mature. If a regulator ever rules them out, mandatory retraining becomes more expensive, which strengthens the case for any cheaper path — but we would reposition toward model quality and insurer evidence, and we've thought that through.

Can we reliably erase a person's data without breaking the model?

Not yet, and we say so. A single fact can be suppressed with little change to overall quality, but a pretrained fact came back after light retraining, and a wider check of an edit from our shipping toolchain found it damaged 8 of 100 other facts. For a person, a rebuild of the fine-tune passes the recovery attacks on one small model, but the usability check is undecided, so by our own rule the test has not passed; cheaper methods have not met their bars. Passing it, and repeating it on a larger model, is the first line of the gate. We publish the result either way.

Is this only a European business?

No, as of this month's research. California's statute has covered personal information inside AI systems since January 2025, its regulator has said that makes it deletable, and its new audit rule requires audit findings to rest on checkable evidence rather than management say-so. No enforcement action has tested it, and we will not be the ones to predict one. What we can say is that the same receipt answers both jurisdictions and the copyright question, and that no vendor on either side of the Atlantic sells weight-level deletion today.

Is this a venture-scale company or a feature?

A feature by default; venture-scale only if our receipt becomes the standard evidence format that governance platforms, auditors and insurers consume. That is the bet the \$1.5M would fund, and we would rather state it than have an investor discover it.

If the technology is copyable, what compounds?

The receipt track record, the published-failures reputation, the measurement method, and outside acceptance of the format. Not the cryptography, not the editing.

Do seven consultancies and three design partners become repeatable paid contracts at \$60–150k?

Unproven until the gate says so. The design-partner deal — free work in exchange for a dated public case study — is built to convert, and the interviews test willingness to pay before we scale anything.

11. Contact

Michael Duane Mooring, Co-founder and CEO · mike@divinci.ai · +1 646-346-0194 · <https://divinci.ai>

Where the outside facts come from

The erasure right is GDPR Article 17; the EU AI Act's own Article 17 is about quality-management systems, and its Annex IV contains no erasure clause — two separate laws. The AI Act delay: Regulation (EU) 2026/1744 (the Digital Omnibus), in force 27 July 2026, moving the strict-rule deadlines to December 2027 and August 2028. CNIL guidance on data-subject rights over AI models, 7 February 2025. EDPB Opinion 28/2024 on when a model counts as anonymous. Copyright: GEMA v. OpenAI,

Munich Regional Court I, 11 November 2025, 42 O 14139/24; GEMA v. Suno, Munich Regional Court I, 31 July 2026. California: AB 1008, chaptered 28 September 2024, effective 1 January 2025, amending Civil Code § 1798.140(v); CPPA staff memorandum to the Board recommending support for AB 1008, 11 July 2024; California Attorney General, Legal Advisory on the Application of Existing California Laws to Artificial Intelligence, 13 January 2025; CPPA regulations on cybersecurity audits, risk assessments and automated decision-making, approved 22 September 2025 (§ 7122(d) on audit evidence); enforcement settlements from the AG and CPPA press releases, 2022–2026 (Sephora \$1.2M, 2022; General Motors \$12.75M, 2026). The Delete Act's deletion platform (DROP) and its \$200-per-request-per-day exposure: CPPA regulations and enforcement advisory, 2025–2026. Transcend and BigID scope statements from their own product pages, September 2026. Our fact-level deletion: divinci.ai/blog/deleting-paris-from-a-language-model, 25 April 2026. Hirundo: seed round announced 9 June 2025; HPE partnership March 2026; Google DeepMind feature May 2026; independent note on its self-reported results, May 2026. Armilla's Lloyd's-backed AI liability policy, launched 30 April 2025, with limits past \$25 million by January 2026; Munich Re and HSB have made comparable moves. Gartner's first Magic Quadrant for AI Governance Platforms, 16 June 2026. All money figures about Divinci itself are our own records.