

Divinci — Making AI Models Forget

Model erasure with a receipt an auditor can verify without us. One page for investors & advisors · 29 Aug 2026

THE PROBLEM

The EU AI Act's enforcement powers went live **2 August 2026**. GDPR's right to erasure was already enforceable — fines up to **€20M or 4% of global turnover**. When a model has *memorised* a person, “we deleted the row” is no longer a complete answer — the person is still in the model. Regulators (EDPB, CNIL) currently treat **retraining** as the defensible response — expensive, slow, and in practice rarely done. Most large enterprises running AI in Europe have **no answer at all** if a model-level erasure demand arrives.

WHAT WE DO

Find what a model has learned about a person. **Remove** it. **Prove** the rest of the model still performs. Hand the auditor a **signed receipt they can check without trusting us** — verifiable with an open-source tool that has no dependency on Divinci, and live in production today. The third step is the one nobody else sells: an unlearning vendor hands back a model, not evidence the model still works.

WHERE WE ARE, PRECISELY

The receipt is in production. The erasure engine runs in staging, on open-weight models. Our end-to-end test in August was **corpus-level erasure of published documents — not model weights, and not personal data**. Whether the technique transfers to an individual's data in weights is exactly what the next 30 days test. No regulator has yet adjudicated whether *any* unlearning method discharges Article 17 — the largest open question in this market, and it is open for everyone in it.

WHAT IT'S WORTH

The substitute for an erasure is **retraining the model: \$50k–500k per event**, repeated per request — the number our pricing anchors to. The qualified market (open-weight, self-hosted, EU-facing models with a live erasure exposure) is **unmeasured**; measuring it is a task in the 30-day gate, not an assumption in this document.

THE PLAN

A **30-day go/no-go gate**, all runnable now: run the individual-level erasure experiment with performance gates on both sides; take a scoped EU counsel opinion on whether receipt-backed suppression discharges Article 17; interview 15–20 privacy counsel and DPO services to size the demand; sign three design partners; publish the evaluation protocol — including our failures.

Direct and design-partner first. The channel through OneTrust, the Big 4 and the governance platforms is the year-two motion; we are not spending a partner introduction before the gate produces something their diligence survives.

Year-1 cost: ~\$224k at founder scale, funded from operations — no raise needed to run the gate. Programme revenue in year 1 is **\$0**: twelve months of investment. The real constraint is focus, not cash — roughly 15% of a three-person team's year, taken out of the delivery work that funds the company. The ~\$1.5M version is only justified *against* the gate's result.

WHY US

- **The receipt is real and shipped.** Ed25519-signed, served from a public endpoint, checked by an MIT-licensed verifier with zero Divinci dependency. Not a format we propose — one that runs.
- **The proof step already caught real damage.** In our own August test, the control topics degraded *more* than the topic we erased (–0.058 vs –0.026). A vendor who cannot measure that ships the damage silently.
- **We published that result ourselves** — including the headline that hurts: deleting a document removed the organisation's viewpoint but not the information; the assistant kept answering the erased question fluently from the model's own parameters. That is the strongest argument for weight-level erasure we own, and it came from our own failed test.
- **What that is and isn't.** Publishing failures is not a moat — a competitor can copy it in one press release. The defensible position is the bundle: baseline, erase, re-measure, attest, as one verifiable artifact. Nobody else sells the measurement half.

WHAT WE'RE ASKING

No capital — the gate runs on operations. This month we need: **three design partners** with open-weight models and EU exposure; an introduction to **EU data-protection counsel** for the scoped question; and a **named academic** willing to adversarially test the result and co-sign it if it holds. Warm introductions to the governance platforms and the Big 4 are worth more in ninety days than today — the trigger is a passing experiment plus the counsel memo.

And the same scrutiny you would give any claim here: every number in our documents is audited against our own code, and where it failed the audit, we say so.

Full strategy document (v0.3, 9 pp.) and detailed plan with line-item expenses, scenarios, pricing and partner tiers (workbook v0.4) available on request. Prepared with Claude Opus 5; every figure audited against the repository.