EU AI Act Compliance Channel Strategy

Getting Divinci's vIndex + LarQL erasure technology into the hands of the compliance vendors who sell to the Fortune 500

v0.2 draft for fleet mastermind — prepared by Hermes Local, 2026-08-28. Anchored to the shipped product at divinci.ai/compliance. Every market figure carries its source; company capability claims are quoted from the product page and marked as company-published. Expense model is a companion workbook (EU-AI-Act-Model-Erasure-Expenses-v0.1.xlsx). Companion decision: after the team settles this draft, run it past Claude Opus 5 via Claude Code.

1. Executive summary

- Enforcement is live. EU AI Act enforcement powers started 2 Aug 2026; the AI Office signaled pre-built cases at launch and first fines were reported within weeks. GDPR erasure rights are already enforceable and fines reach €20M / 4% of global turnover.
- Fortune 500 companies are buying AI compliance at scale: AI governance platform spend is $492M in 2026 heading past $1B by 2030 (Gartner, Feb 2026); EU AI Act compliance software alone is a $1.14B (2026) market growing 28%/yr to $4.05B by 2031 (Mordor Intelligence).
- The one module none of those vendors can sell today is provable erasure of data from a model's weights. Every major vendor has an EU AI Act solution page; none has a model-erasure product with verifiable receipts. That empty slot is visible to their buyers — which is exactly the slot we fill.
- Strategy: do not sell to the Fortune 500 directly. Sell through the vendors who already own them. OneTrust alone reaches over half of the Fortune 500; the Big 4 consultancies run the conformity-assessment work; the GDPR request-handlers (law firms, DPO services) receive the erasure demands.
- Two-prong distribution spine (fleet consensus): Prong A — GDPR Art 17 erasure-with-proof sells to privacy/DPO/GRC buyers (OneTrust DSR, TrustArc, BigID, Securiti, law firms, DPO-as-a-service). Prong B — EU AI Act Annex IV documentation-from-vIndex sells to AI-governance/model-risk buyers (OneTrust AI Gov, IBM watsonx.governance, Credo AI, Holistic AI, ServiceNow, SAP, MetricStream, Enzai). Two buyer sets, two vendor ecosystems, one artifact: the receipt.
- Wedge: certification-grade evidence — and it already exists as a shipped product. vIndex (an interpretable map of a model's learned associations) + LarQL (a patch language over it) produce a portable, SHA-256-checksummed receipt: apply to suppress a fact, remove to restore the model bit-for-bit; the patch IS the audit log. That artifact is exactly what the industry cannot currently produce: no shared evaluation protocol certifies weights-level removal (arXiv:2507.19894 survey; EDPB Jan 2025 report; CNIL guidance), and every competitor's claim is self-measured. A product that erases AND issues a verifiable receipt is the moat.
- This document: regulatory hook (with a scope clarification on 'Annex IV Article 17'), market, vendor landscape, honest competitive check (Hirundo exists), the shipped vIndex + LarQL wedge, five channel tiers, deal structure, 90-day plan, and the expense model.

2. The regulatory hook — and a scope clarification

Scope clarification up front, because this is the foundation of every pitch: there is no single 'Annex IV Article 17' provision. The phrase fuses two distinct obligations — and getting it right stops us pitching a citation that doesn't exist. The erasure obligation lives in GDPR Article 17 (right to erasure). 'Article 17' in the EU AI Act is the quality-management-system article, a different thing entirely. The AI Act's Annex IV (technical documentation) is the documentation/evidence hook. The erasure obligation lives in GDPR Article 17. The

AI Act provisions are the documentation/evidence hooks that make erasure capability sellable. One product, two regulatory stories — and Divinci's compliance page already maps the same requirements (Annex IV, GDPR Art 17, HIPAA, NIST AI RMF, SOC 2, ISO 42001), which validates this reading.

2.1 AI Act Annex IV — Technical documentation (Art 11)

- Providers of high-risk AI systems must draw up technical documentation before placing the system on the market and keep it current for 10 years (Arts 11, 18). Nine sections, per the Annex IV text.
- Section 2(d) is the hook: data requirements — datasheets describing training methodologies, training datasets, provenance, how data was obtained/selected, and data cleaning methodologies.
- Section 2(g): validation and testing procedures, metrics, test logs and signed test reports — the evidence format our receipt artifact is built to feed.
- Breach of provider obligations (incl. failing to maintain Annex IV documentation): fines up to €15M or 3% of worldwide annual turnover (Art 99(4)).

2.2 AI Act Article 17 — Quality Management System

- High-risk providers must run a documented QMS covering 13 mandatory elements — including data management procedures (acquisition, collection, labelling, storage, retention) and record-keeping.
- Timing was moved by the Digital Omnibus (Regulation (EU) 2026/1744): Annex III standalone high-risk obligations apply from 2 December 2027; Annex I product-route high-risk from 2 August 2028. Enforcement powers themselves began 2 Aug 2026.
- EN 18286 ('AI: Quality Management System for EU AI Act Regulatory Purposes') was approved by CEN/CENELEC in June 2026 — the first harmonised standard supporting the Act; OJ citation pending. Providers adopting it get presumption of conformity.
- Cost anchor: QMS setup runs €190k–330k with ~€71.4k/yr maintenance (CEPS cost study) — the budget context a €X-thousand module sells into.

2.3 GDPR Article 17 — Right to erasure (the actual erasure driver)

- Data subjects can demand erasure of their personal data 'without undue delay' — practically one month, extendable to three. GDPR fines reach €20M or 4% of global turnover.
- EDPB (Opinion 28/2024): models from which personal data can be extracted are not anonymous — the GDPR applies to the model itself. EDPB's Jan 2025 report 'Effective implementation of data subjects' rights' walks through retraining, exact unlearning (SISA-style), and the evidence problem — it is the EU authority's own map of our market.
- CNIL guidance: retraining is the defensible route today; output filters as fallback; the state of the art is acknowledged as insufficient. That acknowledged gap is our opening — a receipt-backed patch is a stronger answer than either.
- This is the demand generator: every EU-facing F500 AI deployment is exposed to Art 17 requests it currently cannot answer at the model level. 'We deleted the row' is not a complete answer when the model has memorised the person.

2.4 Enforcement timeline (urgency)

- 2 Aug 2026 — enforcement powers live: prohibitions, GPAI duties, Art 50 transparency; AI Office can send information requests, run model evaluations, require measures, and fine GPAI providers up to €15M / 3% of turnover. Responding slowly to an information request is itself finable.

- Aug 2026 — first fines reported within weeks of launch (one tracker: €47M across hiring-AI, credit-scoring and emotion-recognition cases, 17 Aug 2026; other trackers show a quieter ledger — the point stands that enforcement is no longer theoretical).
- 2 Dec 2026 — end of the Art 50 transparency-marking grace period.
- 2 Dec 2027 / 2 Aug 2028 — high-risk Annex III / Annex I obligations apply.
- Buying implication: the 2026–2027 window is when compliance vendors assemble their AI governance suites and F500 budgets allocate. A module they can launch with this cycle is worth far more than one added after the deadline panic.

3. Market size & timing

Metric	Value	Source
AI governance platform spending	\$492M (2026) → \$1B+ (2030); fragmented AI regulation to cover 75% of world economies by 2030; governance tech cuts regulatory expense ~20%	Gartner press release, Feb 2026
EU AI Act compliance software	\$0.92B (2025) → \$1.14B (2026) → \$4.05B (2031), CAGR 28.06%; software = 72.4% of 2025 spend; services fastest at 30.7% CAGR	Mordor Intelligence, Jul 2026
Machine unlearning solutions (our niche)	\$0.15B (2025) → \$2.73B (2032), CAGR 51.2%	Stratistics MRC
Enterprise buyer spend	Large-enterprise initial compliance \$8–15M; GPAI providers \$12–25M first year; QMS €190–330k setup + €71.4k/yr	CEPS via Axis Intelligence synthesis, 2026
AI TRiSM market (adjacent)	\$2.34B (2024) → \$7.44B (2030), CAGR 21.6%	Industry market analyses, 2024–25

Reading: the platform market is where the revenue sits; the unlearning niche is where we sit; the compliance-spend figures are what our pricing should reference. F500 buyers already budget mid-seven figures for AI compliance — a module priced at tens of thousands per model per year is an easy incremental line item, especially when it de-risks a 4%-of-turnover GDPR exposure.

4. Who sells compliance to the Fortune 500

The landscape splits into four camps, and each is a potential channel rather than a competitor — none of them can build receipt-grade erasure in-house in this cycle.

Vendor	What they sell	F500 traction / notes	Pricing posture	Channel relevance
OneTrust	Privacy/GRC + AI Governance; EU AI Act solution; DSR automation	Over half the Fortune 500; 14,000+ customers; 7 of top-10 healthcare & retail; customers incl. BT, Kuehne+Nagel; partners Databricks,	Custom, enterprise	TIER 1 — the single fastest route to F500 distribution; erasure-with-receipt slot sits next to their DSR module

		Microsoft, Snowflake; Tech Partner Program with Integrations Gallery; KPMG Trusted AI on OneTrust precedent		
IBM watsonx.governance	AI governance + GRC (OpenPages); 200+ frameworks; compliance accelerators	Gartner Magic Quadrant Leader; Zurich Insurance; public pricing ($0.60/resource unit Standard)	Premium / metered	TIER 1 — gravitas + F500 financial-services footprint
Credo AI	Purpose-built AI governance; EU AI Act/NIST/ISO 42001 policy packs; GAIA assistant	Forrester Wave Leader; 30+ hyperscaler/SI/enterprise partners incl. Databricks & Microsoft	Mid-five-figure annual, custom	TIER 1 — best-of-breed governance layer; wants differentiated technical modules
Holistic AI	Audit-led governance; 40+ technical tests; runtime enforcement	London; enterprise; audit heritage	Custom	TIER 1/2 — technical-testing DNA matches receipt evidence
Saidot	EU-native governance knowledge graph	EU market; collaborative workflows	Custom	TIER 2 — EU credibility
Vanta / Drata / Secureframe / Scrut / Modulos	Compliance automation with AI governance modules (ISO 42001, NIST AI RMF)	Mid-market + F500 tech-forward; fast-moving	Annual SaaS	TIER 2 — add 'model erasure evidence' to their control libraries
Microsoft Purview / ServiceNow AI Control Tower / Collibra / SAP	Enterprise suites folding AI governance into existing platforms	F500 defaults; M365/E5 bundled	Bundled / premium	TIER 4 — marketplace/ecosystem route later; slow cycles
Big 4: Deloitte, PwC, EY, KPMG	AI Act strategy, gap analysis, conformity assessment, CE-marking support, maturity assessments	Own the F500 relationship; KPMG runs CE-marking/conformity work; EY invested $1.4B in AI; PwC runs algorithm-validation services; Deloitte runs inventory/classification frameworks	Advisory fees	TIER 3 — the implementation channel; embed as their technical-erasure workbench
Law firms + DPO-	Privacy counsel,	Receive the Art	Hourly/retainer	TIER 5 —

as-a-service + IAPP community	erasure-request handling, DSR response	17 requests; need a vendor-neutral answer		demand-side pull; refer us in
Hirundo (competitor — see §5)	Machine unlearning platform; weight-level surgical edits	Israel; $8M seed; Google DeepMind-featured; HPE + NVIDIA partner	Enterprise SaaS	Competitor / potential partner — see §5

5. Competitive reality check: machine unlearning

Honest finding, stated plainly: 'nobody else on the market' is not accurate. Machine unlearning is a real, funded, growing niche — and the research surfaced a direct competitor doing weight-level work on open-weights models.

- Hirundo (Binyamina, Israel; founded 2023; $8M seed led by Maverick Ventures; ~18 staff) calls itself 'the first machine unlearning platform'. Weight-level surgical edits on open-weights models (Llama 4, DeepSeek-R1, Qwen, Gemma, Cohere Aya) — no retraining from scratch. Featured by Google DeepMind (hardened Gemma 4 in Gemmaverse); HPE Unleash AI partner; NVIDIA GTC showcase; Interop Tokyo Best of Show. Angles today: security hardening (prompt injection −86%, bias reduction) and US model-risk compliance (SR 11-7).
- Hyperscaler R&D: IBM 'Split, Unlearn, Merge' (SPUNGE, Oct 2024); Google/DeepMind unlearning research incl. its Hirundo case study; Microsoft best-practice guidance on approximate unlearning. Academic consolidation: arXiv:2507.19894 (generative-model unlearning survey).
- Academic lineage and the strawman column: rank-one weight edits of the ROME/MEMIT line are published, well-known methods — the page's competitor column (access logs, prompt guardrails, soft-delete) compares against vendors who are not doing weight-level edits at all, so it cannot rule competition in or out. The real question, answered above, is who ships weight-level unlearning: Hirundo plus hyperscaler R&D — and nobody ships receipt-grade, reversible, audit-log-native erasure with EU compliance evidence.
- The gap that matters — certification. No shared evaluation protocol exists to certify that an unlearning operator actually removed data from weights. Every vendor's '99% unlearned' is measured against its own self-selected protocol. No EU regulator has accepted or rejected unlearning as satisfying GDPR Art 17. EDPB and CNIL both treat retrain-and-attest as the only currently defensible path.
- Implication for us: Hirundo's traction validates the market and educates buyers — good. But Hirundo sells 'make the model safer'; the EU buyer needs 'prove the model forgot, with evidence an auditor accepts'. Divinci's answer is already shipped: a LarQL patch carries a SHA-256 checksum, before/after probe results (feature scores, perplexity delta), reversibility — the patch IS the audit log, mappable to Annex IV 2(g) signed test reports. That is exactly what the research says no one else can produce credibly. The strategy is to make that receipt the compliance industry's standard artifact.

6. Divinci's wedge & positioning — the shipped product

The wedge is not hypothetical. Per the company product page (divinci.ai/compliance, company-published claims):

- The product: vIndex maps a model's learned associations (entity → relation → target, with weight, layer, feature); LarQL is the query/patch language over it. A LarQL patch is a portable JSON file with a SHA-256 checksum — apply it to suppress a fact, remove it to restore the model bit-for-bit.

The patch IS the audit log: operators, regulators, and downstream consumers can verify the same operation independently.

- Real shipped example (their public test suite): the Gate-3 patch 'gdpr-art17-paris-capital' on google/gemma-4-E2B-it — a delete op on the Paris→capital association (layer 27, feature 11179): before = feature score 18.10, rank #1; after = feature absent from top-25; Δ perplexity (WikiText-103) +0.02%; vIndex sha256 recorded. Measurable, repeatable, reversible — the receipt. Honest caveat from the page: this demo erases a wrong fact (factual error-correction). GDPR Art 17 sales needs an individual-level demonstration (canary/membership-inference-style erasure of a person's data) — that is the first thing skeptical counsel will probe, and it belongs in the 90-day plan, not assumed.
- Certification status, stated plainly (company-published, from the page): the vIndex dossier is a draft for your assessment, not a declaration of conformity; SOC 2 — not certified; ISO/IEC 42001 — certification not yet held; twelve vIndexes published. Compliance vendors resell certified artifacts to F500 buyers, so certification is a gating workstream for Tier-1 talks — the costs are already in the expense workbook (SOC 2 $35–75k, ISO 42001 $45–110k). The plan must not pretend otherwise.
- Coverage: twelve published vIndexes across Gemma, Qwen, Llama, Mistral, OpenAI MoE, and Microsoft 1-bit — 'bring your model, we'll build the receipt.' Open-weights only, by design: the only models where bit-for-bit reversible erasure is even possible, and the fastest-growing segment of enterprise inference.
- Positioning (their own line): 'Compliance you can prove. On any open transformer.' Channel wording we should use with vendors: 'The last mile of EU AI compliance — erasure with a receipt.'
- Two regulatory stories: (1) GDPR DSR response — legal teams/DPOs: answer Art 17 erasure requests at the model level within the one-to-three-month clock, receipt included; sells against the 4%-of-turnover fine. (2) AI Act technical file — compliance teams: Annex IV technical documentation auto-generated from the vIndex; Annex IV 2(d) datasheets, Art 10 data governance, and the Art 17 QMS data-management element become executable and documentable; sells into the €190–330k QMS budget.
- What must be built for the channel play is NOT the evidence artifact (it exists and ships) — it is the vendor-facing layer: an erasure API + SDK, receipt/report generation in partner formats (workflow-native, Annex IV 2(g)-style signed test logs), and a partner dashboard. That is a platform-engineering sprint, not a research problem — the moat holds because the science is done and the integration surface is ours to define first.

7. Fanned-out go-to-market — five tiers

The strategy is deliberately broad: five channel tiers running in parallel, sequenced by time-to-revenue. Each tier is a distinct motion with its own deal structure.

Tier	Target	Mechanics	Deal structure	Priority
1. AI governance platforms	OneTrust (Tech Partner Program → Integrations Gallery; DSR + AI Governance modules), IBM watsonx.governance, Credo AI, Holistic AI	REST API + SDK; white-label 'model erasure with receipt' module; joint GTM and joint webinar; prong mapping — pitch erasure to their DSR/privacy	Rev-share 15–30% or per-erasure pricing; they get a marquee module, we get F500 distribution	WEEK 1 — fastest route to half the F500

		module owners, documentation to their AI-governance module owners		
2. Compliance automation	Vanta, Drata, Secureframe, Scrut, Modulos	Add 'model erasure evidence' to their AI control libraries; ISO 42001/NIST mapping	Per-control-library license; rev-share	WEEK 4
3. Big 4 + SI alliances	Deloitte, PwC, EY, KPMG; large consultancies	Embed as the technical-erasure workbench inside their AI Act / conformity-assessment offerings; certify their staff; white-label	Implementation fees + per-deployment license; co-branded evidence reports	WEEK 6 — longest cycle, largest deals
4. Marketplaces & MLOps + deployment stack	AWS Marketplace, Azure Marketplace, Google Cloud; Databricks, Dataiku, DataRobot; Hugging Face (vIndexes already there); the serving/inference layer where open-weights actually run — vLLM, Together, Fireworks	Listed product + one-click trial; integration with model registries; open-weights audience natively	Marketplace listing fees + consumption pricing	WEEK 8 — low touch, fills the funnel
5. Request-side ecosystem	Privacy/ employment law firms, DPO-as-a-service, IAPP community, TrustArc/BigID/Securiti DSR platforms	The people who receive Art 17 requests; vendor-neutral referral engine; co-authored guidance	Referral/affiliate + their billable hours	WEEK 2 — creates pull demand
Cross-cut	Direct enterprise design-partner program: 2–3 EU-facing F500 tech companies	Pilot at cost or below in exchange for public case study + receipt validation in a real DSR workflow	First-year pilot pricing; land for case study	WEEK 1 — proof asset for all tiers

Sequencing logic: Tier 1 gets revenue and legitimacy, Tier 5 generates the demand pull ('our counsel needs this'), Tier 3 delivers the whale deals, Tiers 2/4 fill volume. The design-partner program funds the receipts that every tier sells with.

8. Deal structure & pricing levers

- Event pricing: per-erasure run with receipt (e.g., per GDPR request or per dataset-withdrawal) + per-model annual coverage fee for ongoing erasure-readiness.
- Platform license: annual module license priced per governed model or per org — positioned against the F500's $8–15M compliance budgets and the €190–330k QMS spend; a tens-of-thousands-per-year module is an easy increment.
- OEM/white-label: 15–30% rev-share to Divinci when a Tier-1/3 partner sells under their brand.
- Services: receipt/report generation, integration, certification support, and joint training — services are the fastest-growing segment of the market (30.7% CAGR).
- Pricing north star: price against risk, not cost. A module that demonstrably closes a 4%-of-global-turnover GDPR exposure or a €15M/3% AI Act documentation fine carries its own ROI story.

9. 90-day execution plan

1. Days 0–30: build the vendor-facing layer (erasure API + SDK + receipt/report generation); build the individual-level erasure demo (canary/membership-inference style) that answers the demo-category objection; kick off SOC 2 / ISO 42001 certification program (gating for Tier-1); sign design partner #1; open the OneTrust Technology Partner Program application; brief two privacy law firms (Tier 5).
2. Days 31–60: Tier-1 proposals out (OneTrust, IBM watsonx.governance, Credo AI); AWS/Azure marketplace listings in flight; IAPP presence; design partner #2 in the EU.
3. Days 61–90: three signed LOIs/MOUs across at least two tiers; first joint webinar with a Tier-1 or Tier-3 partner; reconcile spend against the expense model; then run the whole plan past Claude Opus 5 (Claude Code) for a second opinion before committing year-2 budgets.

10. Expenses

Companion workbook: EU-AI-Act-Model-Erasure-Expenses-v0.1.xlsx. Three scenarios for the first 12 months (lean validation / standard land-3-5-partners / aggressive fanned-out), with line items across engineering, security & certification (SOC 2, ISO 42001), legal & regulatory advisory, GTM, partner programs, marketing, compute, services absorbed on design partners, travel, overhead, and contingency. All figures are planning estimates with stated bases; nothing is quoted vendor pricing unless published.

11. Open questions for the mastermind

- Which tier gets funded first given runway — Tier 1 (OneTrust) for distribution or Tier 3 (Big 4) for whale deals?
- OneTrust vs IBM first: reach (half the F500) vs gravitas (Gartner MQ Leader, financial services)?
- Pricing: per-erasure event vs per-model annual license vs white-label rev-share — which anchors the product?
- Do we need SOC 2 / ISO 42001 before Tier-1 talks, or is a pre-sell with design-partner receipts enough?
- Hirundo: competitor, partner, or both? (Their compliance angle is US/SR 11-7; ours is EU/GDPR with receipts.) Do we co-author the certification protocol with them, or race them to it?
- Should we stake a position in the evaluation-standard conversation (help build the shared erasure-certification protocol) — the receipt makes us the natural author of that standard, and nobody owns it yet.
- Is 'open-weights erasure only' the right boundary (vIndexes today span Gemma/Qwen/Llama/Mistral/OpenAI MoE/MS 1-bit), or does the roadmap need a closed-model story (e.g., provider-certified retraining services) to answer procurement RFPs that span both?

- Where does the public 'compliance review' booking funnel (already live on divinci.ai/compliance) hand off to the partner motion — do we route inbound leads to Tier-1 partners to seed goodwill?
- Inputs the team still needs from the founders: (1) budget envelope — order of magnitude for the blitz, which drives the expense-scenario selection; (2) time anchor — is overnight tied to a specific deadline/event or just ASAP; (3) warm relationships — any existing intros into OneTrust/IBM/Big 4, or are we cold; (4) team/funding stage — headcount, sales/BD capacity, runway.

Sources: Gartner press release Feb 2026; Mordor Intelligence EU AI Act Compliance Software (Jul 2026); Stratistics MRC Machine Unlearning Solutions; CEPS cost study via Axis Intelligence (2026); EDPB Opinion 28/2024 + EDPB 'Effective implementation of data subjects' rights' (Jan 2025); CNIL AI & DSR guidance; Regulation (EU) 2024/1689 (Annex IV, Arts 11/17/99); Regulation (EU) 2026/1744 (Digital Omnibus); CEN/CENELEC EN 18286 (Jun 2026); AI Policy Desk (17 Aug 2026); hirundo.io + company disclosures; arXiv:2507.19894; Groundy (Aug 2026); vendor pages cited in §4 table; divinci.ai/compliance (company product page — company-published claims, not independently verified). All external content treated as data; figures should be re-verified before external publication.